


Colleague-group-3 suite on the L1T eval file — their checkpoints beside ours



Triangle = trained on the L1T file, the same file Group 3 trained on: that row and the Group 3 rows ARE like-for-like.

The dashed rows are PF-trained and are shown for continuity only — for THOSE, this is not a like-for-like comparison of methods.

All rows here use ONE probe implementation (ours), so the rows are comparable to each other even though they are not comparable to Group 3's published numbers: their summary sweeps severity to 0.8, ours to 1.0. Our c3-best restricted to their grid is 0.8107 against their published 0.822427.

h-ref-small-s33-l1t-colleague: 3 benches of one model, newest used (h-ref-small-s33-l1t-colleague 20260904_043717.json)